

Medium

 Search Write

★ Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



Implementation of Diffusion Model

Hiroaki Kubo · [Follow](#)

4 min read · Dec 14, 2024



I wrote this article because I implemented diffusion model from scratch using the loss derived in [previous article](#).

In this case, I used the swiss roll dataset and the algorithm basically followed [Denoising Diffusion Probabilistic Models](#). The code which I implemented is [here](#).

...

Overview

When we implement the Diffusion model, the following parts need to be implemented. I will explain each of these steps.

- Forward process
- Neural network for training
- Training
- Sampling

• • •

Forward process

The forward process of adding Gaussian noise to the input data is expressed by the following equation. Diffusion model gradually adds Gaussian noise to the data according to a variance schedule β_1, \dots, β_T .

$$q(x_{1:T}|x_0) := \prod_{t=1}^T q(x_t|x_{t-1}), \quad q(x_t|x_{t-1}) := \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

A notable property of the forward process is that it admits sampling x_t , at an arbitrary timestep t in closed form: using the notation $\alpha_t := 1 - \beta_t$ and $\alpha_t^- = \prod_{s=1}^t \alpha_s$, we have

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\alpha_t^-}x_0, (1 - \alpha_t^-)I)$$

Thus, we can get x_t by using a reparameterization trick.

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon \quad (\epsilon \sim \mathcal{N}(0, I))$$

The code to get α_t is here.

```
def calculate_parameters(diffusion_steps, min_beta, max_beta):
    step = (max_beta - min_beta) / diffusion_steps
    beta_ts = torch.arange(min_beta, max_beta + step, step)

    alpha_ts = 1 - beta_ts
    bar_alpha_ts = torch.cumprod(alpha_ts, dim=0)

    return beta_ts, alpha_ts, bar_alpha_ts
```

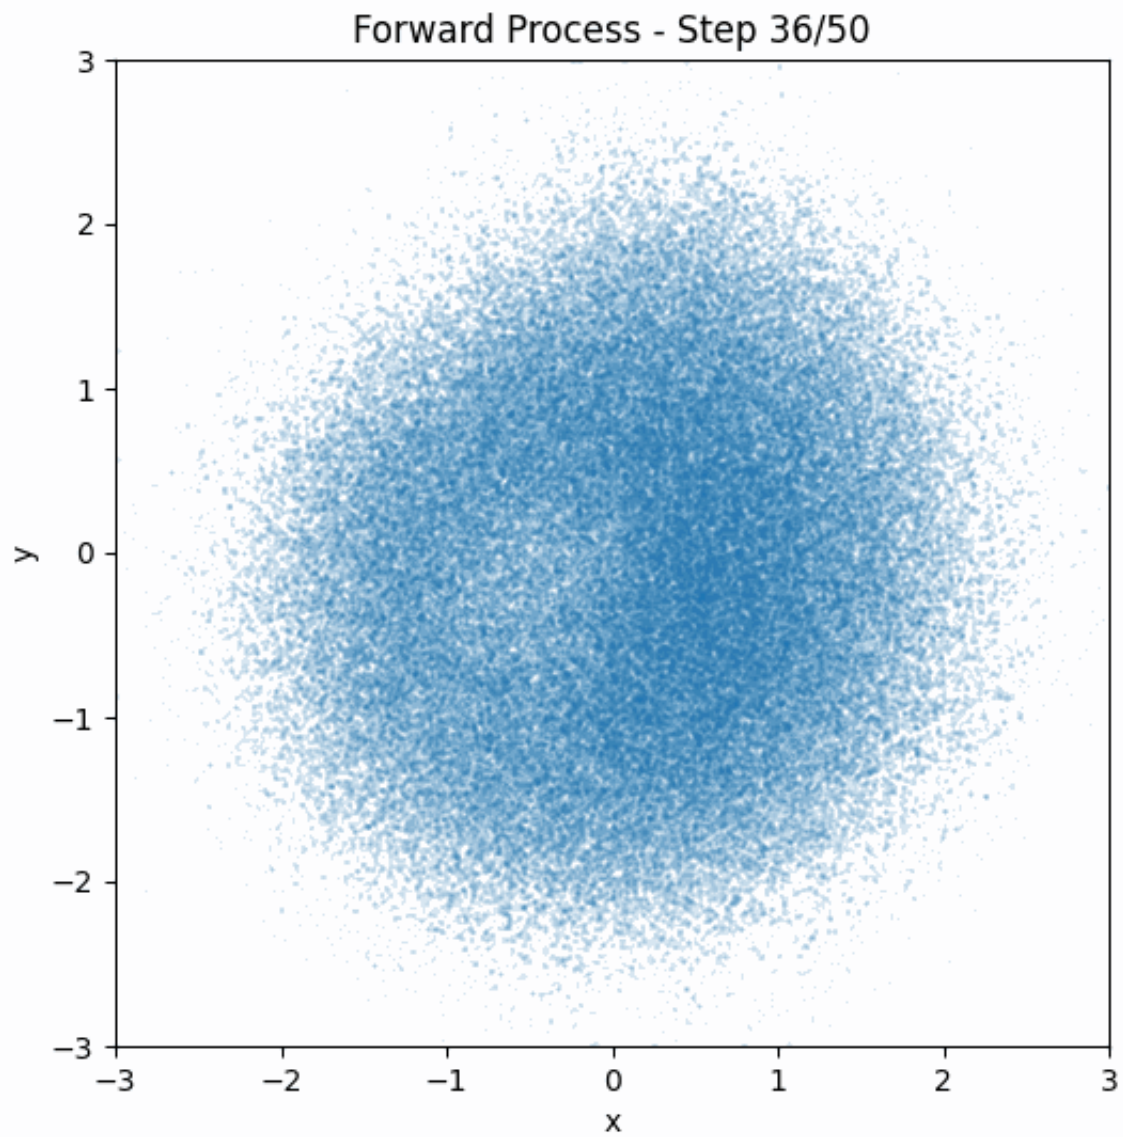
The code to retrieve data and ϵ at any given time in the forward process is as follows.

```
def calculate_data_at_certain_time(x_0, bar_alpha_ts, t):
    eps = torch.randn(size=x_0.shape)
    noised_x_t = (
        torch.sqrt(bar_alpha_ts[t]) * x_0 + torch.sqrt(1 - bar_alpha_ts[t]) *
    )

    return noised_x_t, eps
```

You can try a forward process with swiss roll by running this code. We set the forward process variance constants increasing linearly from β_1

$=0.0004$ to $\beta T=0.02$.



• • •

Neural network for training

Original paper uses U-Net backbone, but I used simple neural network for training this time because it is enough in this data. It has 4 hidden layers and use ReLU as an activation function.

If you want to check the architecture of model, you can run [this code](#).

. . .

Training

In reverse process, we calculate x_{t-1} from x_t and timestep t by normal distribution.

$$x_{t-1} \sim \mathcal{N}(\mu_{\theta}(x_t, t), \Sigma_t)$$

The variance is fixed, so we need to predict only $\mu_{\theta}(x_t, t)$. $\mu_{\theta}(x_t, t)$ can be rewritten as follows by simplifying equations. If you want to know the details of that, please check [my previous article](#).

$$\mu_{\theta}(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right)$$

Therefore, we train ϵ during training, loss function is as follows. If you want to know the details of that, please check [my previous article](#).

$$L_{simple} := E_{t, x_0, \epsilon} \left[\left| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t) \right|^2 \right]$$

I followed below training algorithms of original paper.

Algorithm 1 Training

- 1: **repeat**
 - 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
 $\nabla_{\theta} \left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t) \right\|^2$
 - 6: **until** converged
-

Parameters which I used during training is as follows.

- Optimizer -> Adam
- Batch size -> 128
- Epochs -> 30
- Diffusion timesteps -> 50
- Minimum beta -> 1e-4
- Maximum beta -> 0.02

You can train the diffusion model by running [this code](#).

• • •

Sampling

To sample $x_{t-1} \sim p_\theta(x_{t-1}|x_t)$ is to compute below equation. It is used reparameterization trick.

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z, \quad (z \sim \mathcal{N}(0, I))$$

I followed below sampling algorithms of original paper. We initialize x_T with a random value, and with x_t and t as inputs, the neural network output ϵ , so we use that value to calculate x_{t-1} . This is repeated until x_0 is calculated.

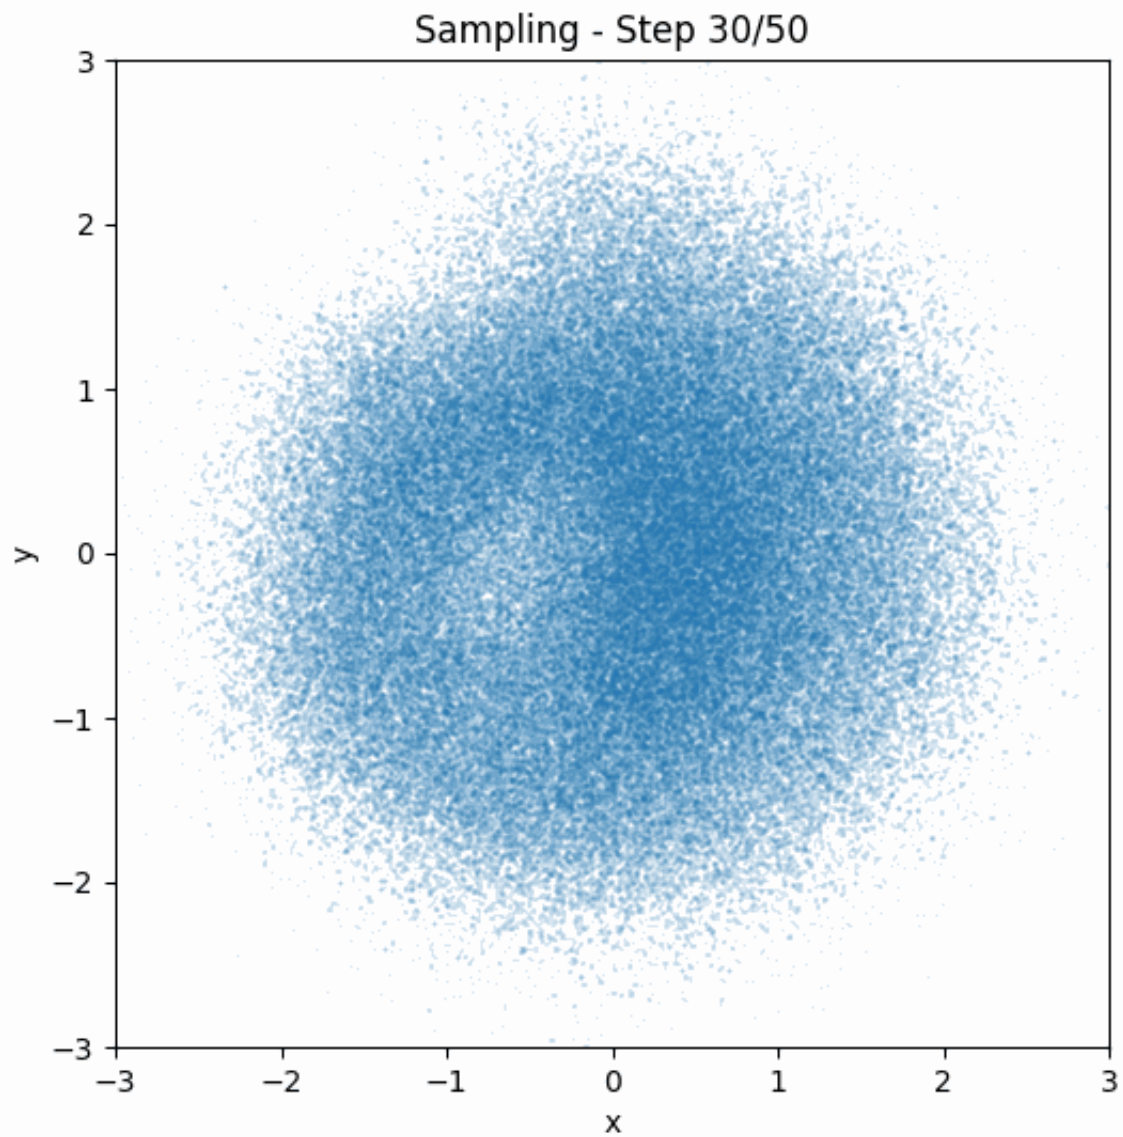
Algorithm 2 Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 

```

You can try sampling by running [this code](#).



• • •

References

- [Denoising Diffusion Probabilistic Models](#)
- [Deep Unsupervised Learning using Nonequilibrium Thermodynamics](#)
- [The Reparameterization Trick](#)
- [toy-diffusion](#)
- [Understanding the Diffusion Model](#)

AI

Python

Diffusion Models

Generative Ai Tools

Neural Networks

**Written by Hiroaki Kubo**

140 Followers · 4 Following

[Follow](#)

Computer Vision Engineer. Computer Vision / Machine Learning / Computer Graphics, <https://www.linkedin.com/in/hiroaki-kubo-2819951ba/>